

-----Project-----

.....Heart Disease Prediction.....

Muhammad Arham Adeel

import libraries

```
In [1]: import pandas as pd
import numpy as np

# data visualization
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

from sklearn.model_selection import train_test_split

#model validation
from sklearn.metrics import log_loss, precision_score, f1_score, recall_score, roc_auc_score
from sklearn.metrics import confusion_matrix, accuracy_score, matthews_corrcoef
from sklearn import metrics

# machine learning algorithms
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier, ExtraTree
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
import xgboost as xgb

from scipy import stats
```

```
In [2]: dt = pd.read_csv("heart.csv")
dt.head()
```

```
Out[2]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3	0
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3	0
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3	0
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3	0
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2	0

here i can check the quick information of our data

In [3]:

`dt.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1025 non-null  int64
1   sex         1025 non-null  int64
2   cp          1025 non-null  int64
3   trestbps    1025 non-null  int64
4   chol        1025 non-null  int64
5   fbs         1025 non-null  int64
6   restecg     1025 non-null  int64
7   thalach     1025 non-null  int64
8   exang       1025 non-null  int64
9   oldpeak     1025 non-null  float64
10  slope       1025 non-null  int64
11  ca          1025 non-null  int64
12  thal        1025 non-null  int64
13  target      1025 non-null  int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
```

Checking the null values

In [4]:

`dt.isnull().sum()`

```
Out[4]: age         0
sex         0
cp          0
trestbps    0
chol        0
fbs         0
restecg     0
thalach     0
exang       0
oldpeak     0
slope       0
ca          0
thal        0
target      0
dtype: int64
```

Check the Shape of the data

In [5]:

`dt.shape`

Out[5]: (1025, 14)

In [6]:

```
rows, columns = dt.shape
print("The total rows/instances in our data is: ", rows)
print("The total columns/labels in our data is: ", columns)
```

```
The total rows/instances in our data is: 1025
The total columns/labels in our data is: 14
```

Here's the summary statistics of our data

In [7]:

```
dt.describe()
```

Out[7]:

	age	sex	cp	trestbps	chol	fbs	restecg	t
count	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000	1025.000000
mean	54.434146	0.695610	0.942439	131.611707	246.000000	0.149268	0.529756	149.000000
std	9.072290	0.460373	1.029641	17.516718	51.59251	0.356527	0.527878	23.000000
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000
25%	48.000000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000	132.000000
50%	56.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000	152.000000
75%	61.000000	1.000000	2.000000	140.000000	275.000000	0.000000	1.000000	166.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000



In []:

Now i can separte the numerical columns and categorical columns

In [8]:

```
df_num = ["age", "trestbps", "chol", "thalach", "oldpeak"]
df_cat = ["sex", "cp", "fbs", "restecg", "exang", "slope", "ca", "thal", "target"]
```

here i see the age distribution

In [9]:


```
sns.distplot(x= dt["age"])
```

H:\download\Anaconda\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[9]:

```
<AxesSubplot:ylabel='Density'>
```




here i can see the distribution of resting bp

```
In [10]: sns.distplot(x= dt["trestbps"])
```

H:\download\Anaconda\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

```
Out[10]: <AxesSubplot:ylabel='Density'>
```




here i can see the distribution of cholestrol

```
In [11]: sns.distplot(x= dt["chol"])
```

H:\download\Anaconda\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

```
Out[11]: <AxesSubplot:ylabel='Density'>
```



Checking the uniqueness and counts in multiple columns

```
In [12]: # checking the unique values of the columns
dt["sex"].unique()
```

```
Out[12]: array([1, 0], dtype=int64)
```

```
In [13]: dt["sex"].value_counts()
```


```
Out[13]: 1    713
         0    312
         Name: sex, dtype: int64
```

```
In [14]: sns.countplot(dt["sex"], data = dt)
```

H:\download\Anaconda\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

```
Out[14]: <AxesSubplot:xlabel='sex', ylabel='count'>
```



```
In [15]: # Plotting Continuous/Numerical columns
```

```
c = 1
plt.figure(figsize=(20,40))

for j in df_num:
    plt.subplot(6,3,c)
    sns.distplot(dt[j])
    c = c+1
plt.show()
```

H:\download\Anaconda\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

H:\download\Anaconda\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

H:\download\Anaconda\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

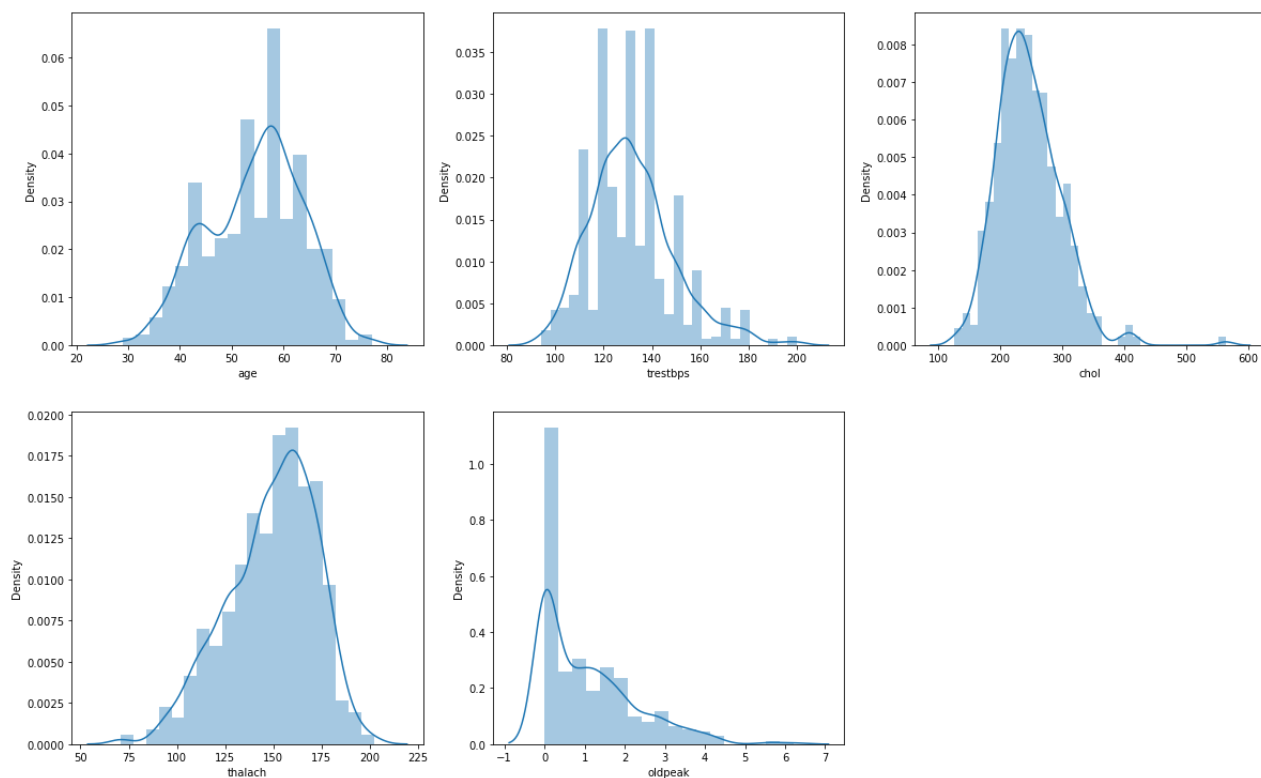
warnings.warn(msg, FutureWarning)

H:\download\Anaconda\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

H:\download\Anaconda\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


warnings.warn(msg, FutureWarning)



Distribution of Heart disease (target variable)

```
In [16]: fig = plt.figure(figsize=(14,6))
ax1 = dt['target'].value_counts().plot.pie( x="Heart disease" ,y='no.of patients',
      autopct = "%1.0f%",labels=["Heart Disease","Normal"]);
ax1.set(title = 'Percentage of Heart disease patients in this dataset')
plt.show()
```

Percentage of Heart disease patients in this dataset



Gender & Agewise Distribution

```
In [17]: plt.figure(figsize=(18,12))
plt.subplot(221)
dt["sex"].value_counts().plot.pie(autopct = "%1.0f%", colors = sns.color_palette("prism"))
plt.title("Gender Wise Distribution")


#here i can check the distribution of age

plt.subplot(222)
ax = sns.distplot(dt['age'])
plt.title("Age wise distribution")
plt.show()
```


H:\download\Anaconda\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

Gender Wise Distribution



Age wise distribution



Distribution of Chest Pain Type

```
In [18]: # plotting normal patients
fig = plt.figure(figsize=(15,5))
ax1 = plt.subplot2grid((1,2),(0,0))
sns.countplot(dt['cp'])
plt.title('CHEST PAIN For NORMAL PATIENTS', fontsize=15, weight='bold')



#plotting heart patients
ax1 = plt.subplot2grid((1,2),(0,1))
sns.countplot(dt['cp'], palette='viridis')
plt.title('CHEST PAIN for HEART PATIENTS', fontsize=15, weight='bold')
plt.show()
```

H:\download\Anaconda\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

H:\download\Anaconda\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will res

ult in an error or misinterpretation.
warnings.warn(



Distribution of Rest ECG

In [19]:

```
# plotting normal patients
fig = plt.figure(figsize=(15,5))
ax1 = plt.subplot2grid((1,2),(0,0))
sns.countplot(dt['restecg'])
plt.title('REST ECG For NORMAL PATIENTS', fontsize=15, weight='bold')


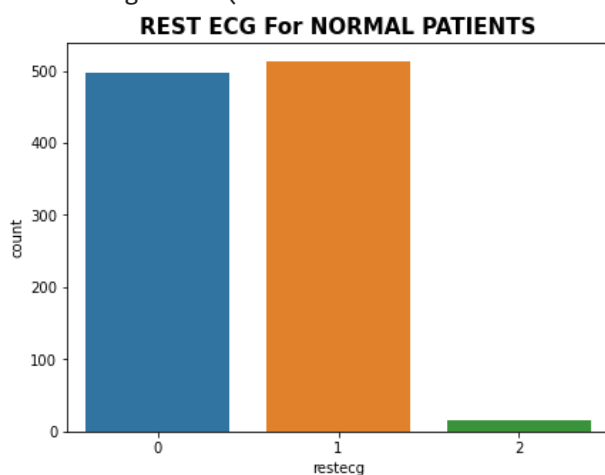
#plotting heart patients
ax1 = plt.subplot2grid((1,2),(0,1))
sns.countplot(dt['restecg'], palette='viridis')
plt.title('REST ECG For HEART PATIENTS', fontsize=15, weight='bold' )
plt.show()
```

H:\download\Anaconda\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

H:\download\Anaconda\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



Distribution of slope

In [20]:

```
# plotting normal patients
fig = plt.figure(figsize=(15,5))
ax1 = plt.subplot2grid((1,2),(0,0))
sns.countplot(dt['slope'])
plt.title('ST SLOPE OF NORMAL PATIENTS', fontsize=15, weight='bold')

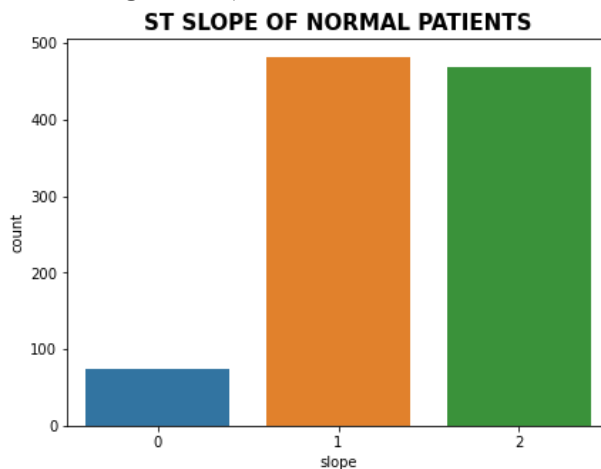
#plotting heart patients
ax1 = plt.subplot2grid((1,2),(0,1))
sns.countplot(dt['slope'], palette='viridis')
plt.title('ST SLOPE OF HEART PATIENTS', fontsize=15, weight='bold' )
plt.show()
```

H:\download\Anaconda\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

H:\download\Anaconda\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

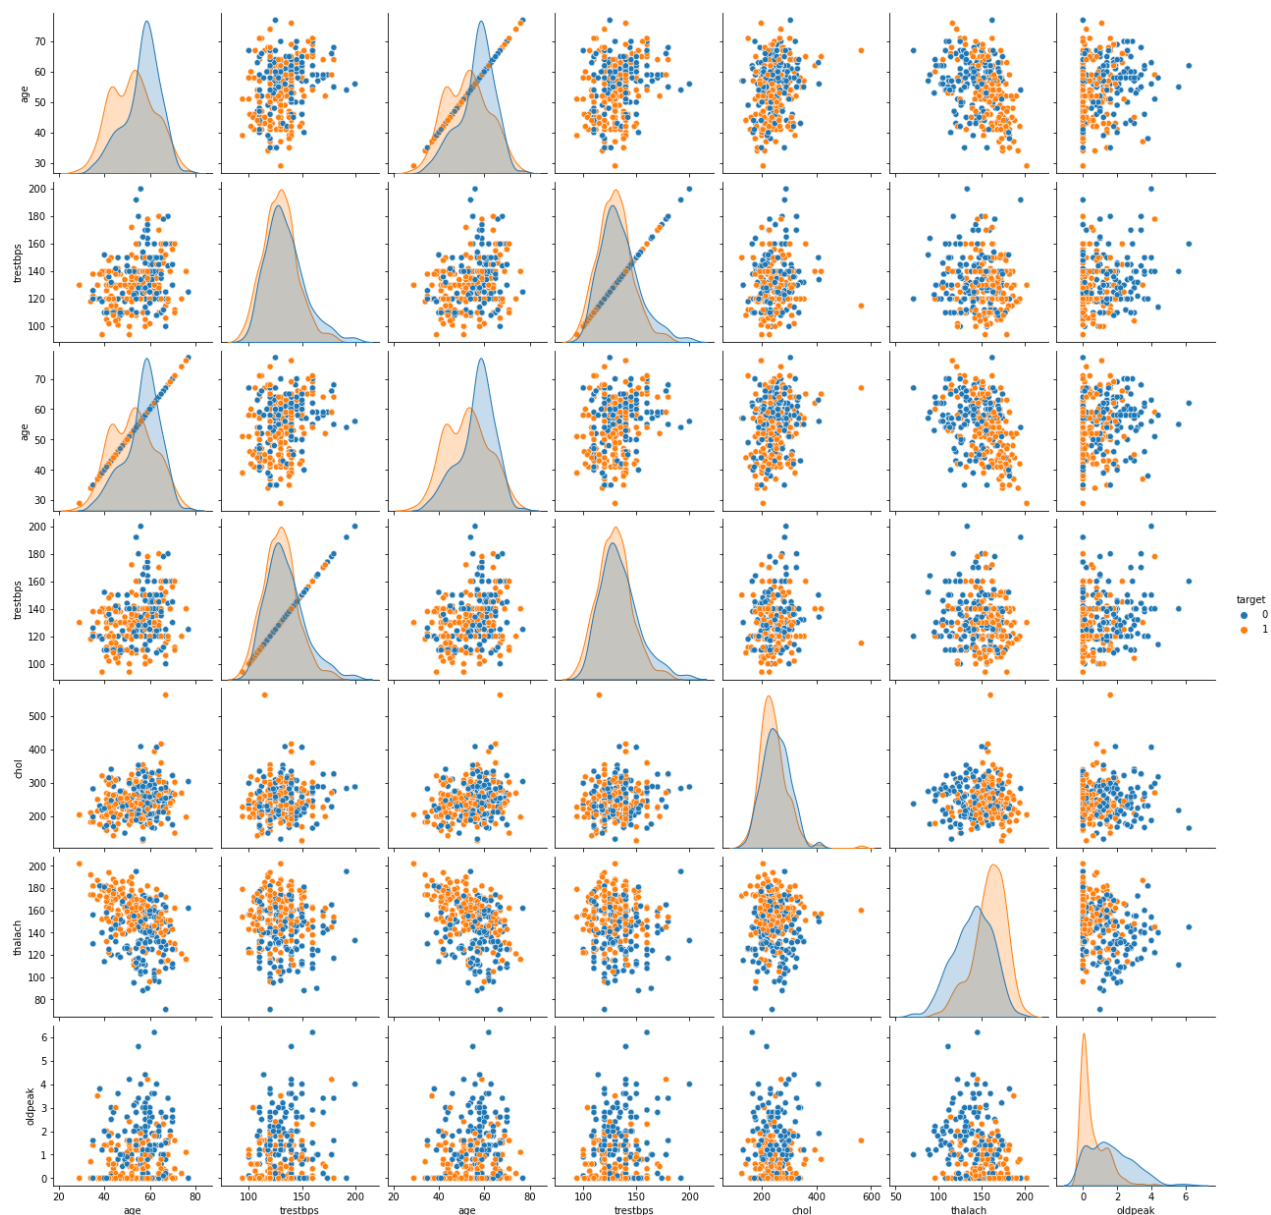


In [21]:

```
sns.pairplot(dt, hue = 'target', vars = ['age', "trestbps", "age", "trestbps", "chol", "thal
```

Out[21]:


```
<seaborn.axisgrid.PairGrid at 0x1bea1d83880>
```



Here i can draw a scatter plot


```
In [22]: sns.scatterplot(x = 'trestbps', y = 'chol', hue = 'target', data = dt)
```

```
Out[22]: <AxesSubplot:xlabel='trestbps', ylabel='chol'>
```



```
In [23]: sns.scatterplot( x = 'trestbps', y = 'age', hue = 'target', data = dt)
```

```
Out[23]: <AxesSubplot:xlabel='trestbps', ylabel='age'>
```



Here i can remove the outlier using

```
In [24]: df_num = dt[["age", "trestbps", "chol", "thalach", "oldpeak"]]
```

```
In [25]: df_num.head()
```

```
Out[25]:
```

	age	trestbps	chol	thalach	oldpeak
0	52	125	212	168	1.0
1	53	140	203	155	3.1
2	70	145	174	125	2.6
3	61	148	203	161	0.0
4	62	138	294	106	1.9

here i can import z score from scipy stats

```
In [26]: # calculating zscore of numeric columns in the dataset

from scipy import stats
z = np.abs(stats.zscore(df_num))
print(z)
```

	age	trestbps	chol	thalach	oldpeak
0	0.268437	0.377636	0.659332	0.821321	0.060888
1	0.158157	0.479107	0.833861	0.255968	1.727137
2	1.716595	0.764688	1.396233	1.048692	1.301417
3	0.724079	0.936037	0.833861	0.516900	0.912329
4	0.834359	0.364875	0.930822	1.874977	0.705408
...
1020	0.503520	0.479107	0.484803	0.647366	0.912329
1021	0.613800	0.377636	0.232705	0.352873	1.471705
1022	0.819834	1.234378	0.562371	1.353113	0.060888
1023	0.488996	1.234378	0.155137	0.429923	0.912329
1024	0.047877	0.663216	1.124743	1.570556	0.279688

[1025 rows x 5 columns]

```
In [27]: # Defining threshold for filtering outliers
threshold = 3
print(np.where(z > 3))
```

```
(array([ 54,  55,  69, 123, 151, 158, 175, 179, 192, 246, 267, 294, 296,
        326, 378, 393, 450, 464, 481, 508, 526, 559, 613, 641, 665, 685,
        688, 833, 889, 958, 996], dtype=int64), array([4, 4, 4, 2, 1, 2, 1, 2, 2, 1, 3,
        1, 3, 1, 3, 4, 2, 2, 2, 1, 4, 3,
        4, 2, 2, 2, 1, 4, 2, 2, 2], dtype=int64))
```

```
In [28]: #filtered where threshold is less than 3
dt = dt[(z < 3).all(axis=1)]
```

```
In [29]: dt.shape
```

```
Out[29]: (994, 14)
```

Correlation

```
In [30]: corr = dt.corr()
corr
```


```
Out[30]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang
age	1.000000	-0.087280	-0.069128	0.285929	0.200024	0.122203	-0.126012	-0.392439	0.102357
sex	-0.087280	1.000000	-0.048426	-0.053918	-0.145201	0.049226	-0.091359	-0.042984	0.145504
cp	-0.069128	-0.048426	1.000000	0.065850	-0.108144	0.077960	0.044571	0.297943	-0.401426
trestbps	0.285929	-0.053918	0.065850	1.000000	0.136774	0.161166	-0.096239	-0.058570	0.050968

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang
chol	0.200024	-0.145201	-0.108144	0.136774	1.000000	0.011746	-0.108515	-0.052669	0.091332
fbs	0.122203	0.049226	0.077960	0.161166	0.011746	1.000000	-0.094645	-0.010288	0.041767
restecg	-0.126012	-0.091359	0.044571	-0.096239	-0.108515	-0.094645	1.000000	0.075493	-0.071684
thalach	-0.392439	-0.042984	0.297943	-0.058570	-0.052669	-0.010288	0.075493	1.000000	-0.390851
exang	0.102357	0.145504	-0.401426	0.050968	0.091332	0.041767	-0.071684	-0.390851	1.000000
oldpeak	0.203698	0.146967	-0.143971	0.139878	0.053062	0.005451	-0.033572	-0.356368	0.320477
slope	-0.163046	-0.052830	0.106634	-0.090646	-0.010255	-0.059806	0.083985	0.385838	-0.257997
ca	0.272587	0.155924	-0.164602	0.065880	0.058614	0.135036	-0.052260	-0.234226	0.112797
thal	0.065701	0.229561	-0.153447	0.022886	0.072217	-0.042327	-0.002379	-0.107926	0.193995
target	-0.229328	-0.297362	0.420497	-0.110103	-0.122853	-0.047673	0.133995	0.425543	-0.442332

In [31]:

```
plt.figure(figsize=(20,16))
sns.heatmap(corr,annot=True,cmap="YlGnBu")
plt.show()
```




Draw a boxplot

```
In [32]: sns.boxplot(dt.age)
```

H:\download\Anaconda\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

```
Out[32]: <AxesSubplot:xlabel='age'>
```




```
In [33]: sns.boxplot(dt["trestbps"])
```

H:\download\Anaconda\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

```
Out[33]: <AxesSubplot:xlabel='trestbps'>
```




```
In [34]: sns.boxplot(dt["cho1"])
```

H:\download\Anaconda\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[34]: <AxesSubplot:xlabel='chol'>




In [35]: `sns.boxplot(dt["thalach"])`

H:\download\Anaconda\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[35]: <AxesSubplot:xlabel='thalach'>




In [36]: `sns.boxplot(dt["thalach"])`

H:\download\Anaconda\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```


Out[36]: <AxesSubplot:xlabel='thalach'>



In [37]: `dt.groupby(by= "sex").sum()`

Out[37]:

	age	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
sex													
0	16199	302	38668	74132	35	179	44075	67	223.9	424	143	610	220
1	37720	649	91571	167911	111	356	104470	267	784.5	970	589	1686	300

In [38]: `dt.groupby(by= ["sex", "cp"]).count()`

Out[38]:

	age	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
sex cp												
0 0	119	119	119	119	119	119	119	119	119	119	119	119
1 57	57	57	57	57	57	57	57	57	57	57	57	57
2 103	103	103	103	103	103	103	103	103	103	103	103	103
3 13	13	13	13	13	13	13	13	13	13	13	13	13
1 0	356	356	356	356	356	356	356	356	356	356	356	356
1 107	107	107	107	107	107	107	107	107	107	107	107	107
2 175	175	175	175	175	175	175	175	175	175	175	175	175
3 64	64	64	64	64	64	64	64	64	64	64	64	64

In [39]: `grp = dt.groupby(by= ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'ca', 'thal', 't
grp.head(20)`

Out[39]:

	age	trestbps	chol	thalach	oldpeak
sex cp fbs restecg exang slope ca thal target					

										age	trestbps	chol	thalach	oldpeak
sex	cp	fbs	restecg	exang	slope	ca	thal	target						
0	0	0	0	0	0	2	2	0	4	4	4	4	4	4
					1	0	2	1	12	12	12	12	12	12
					2	3	0	3	3	3	3	3	3	3
					3	3	0	4	4	4	4	4	4	4
					2	0	2	0	4	4	4	4	4	4
								1	6	6	6	6	6	6
					1	2	1	3	3	3	3	3	3	3
			1	1	0	2	1	6	6	6	6	6	6	6
						3	0	4	4	4	4	4	4	4
		1	0	1	0	2	1	8	8	8	8	8	8	8
					2	2	1	3	3	3	3	3	3	3
					2	0	2	1	11	11	11	11	11	11
					2	2	1	3	3	3	3	3	3	3
				1	1	0	2	0	9	9	9	9	9	9
						3	0	6	6	6	6	6	6	6
					2	2	0	4	4	4	4	4	4	4
					2	0	2	1	6	6	6	6	6	6
		2	1	1	0	2	0	4	4	4	4	4	4	4
					1	3	0	4	4	4	4	4	4	4
	1	0	1	1	0	3	0	4	4	4	4	4	4	4

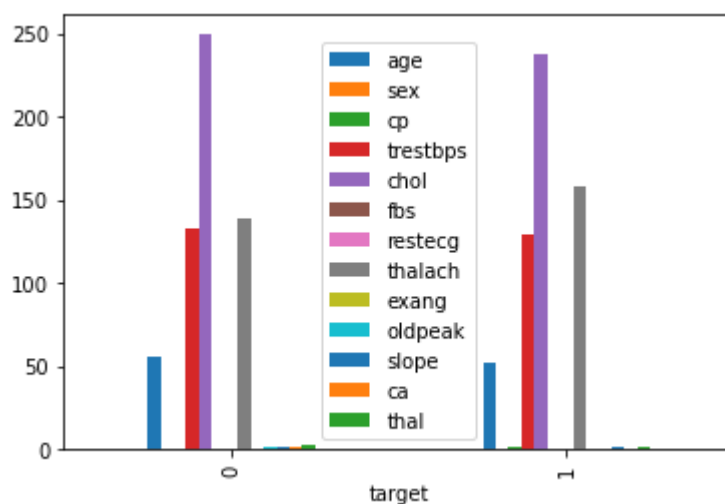
In [40]: `dt.groupby(by= "target").mean()`

Out[40]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang
target									
0	56.430380	0.848101	0.502110	132.955696	249.432489	0.164557	0.464135	139.409283	0.554852
1	52.251923	0.576923	1.371154	129.265385	238.100000	0.130769	0.605769	158.586538	0.136538

In [41]: `dt.groupby(by= "target").mean().plot(kind= 'bar')`

Out[41]: `<AxesSubplot:xlabel='target'>`



checking the unique values of each column

```
In [42]: dt["cp"].unique()
```

```
Out[42]: array([0, 1, 2, 3], dtype=int64)
```

```
In [43]: dt["cp"].value_counts()
```

```
Out[43]: 0    475
         2    278
         1    164
         3     77
         Name: cp, dtype: int64
```

```
In [44]: dt["fbs"].value_counts()
```

```
Out[44]: 0    848
         1    146
         Name: fbs, dtype: int64
```

```
In [45]: dt["restecg"].value_counts()
```

```
Out[45]: 1    505
         0    474
         2     15
         Name: restecg, dtype: int64
```

```
In [46]: dt["exang"].value_counts()
```

```
Out[46]: 0    660
         1    334
         Name: exang, dtype: int64
```

```
In [47]: dt["slope"].value_counts()
```

```
Out[47]: 1    468
         2    463
```

```
0    63
Name: slope, dtype: int64
```

```
In [48]: dt["ca"].value_counts()
```

```
Out[48]: 0    567
         1    220
         2    127
         3     62
         4     18
         Name: ca, dtype: int64
```

```
In [49]: dt["thal"].value_counts()
```

```
Out[49]: 2    537
         3    386
         1     64
         0      7
         Name: thal, dtype: int64
```

```
In [50]: dt["target"].value_counts()
```

```
Out[50]: 1    520
         0    474
         Name: target, dtype: int64
```

```
In [51]: X = dt.drop(columns="target")
         y = dt["target"]
```

```
In [52]: X.head()
```

```
Out[52]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2

```
In [53]: y.head()
```

```
Out[53]: 0    0
         1    0
         2    0
         3    0
         4    0
         Name: target, dtype: int64
```

```
In [54]: print(f"Feature Matrix X shape {X.shape}")
         print(f"Target Vector y shape {y.shape}")
```

Feature Matrix X shape (994, 13)
Target Vector y shape (994,)

Model training and evaluation

```
In [55]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=4
```

```
In [56]: ## checking distribution of target variable in train test split
print('Distribution of target variable in training set')
print(y_train.value_counts())

print('Distribution of target variable in test set')
print(y_test.value_counts())
```

Distribution of target variable in training set

1 412

0 383

Name: target, dtype: int64

Distribution of target variable in test set

1 108

0 91

Name: target, dtype: int64

```
In [57]: print('-----Training Set-----')
print(X_train.shape)
print(y_train.shape)

print('-----Test Set-----')
print(X_test.shape)
print(y_test.shape)
```

-----Training Set-----

(795, 13)

(795,)

-----Test Set-----

(199, 13)

(199,)

```
In [58]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_train[["age", "trestbps", "chol", "thalach", "oldpeak"]] = scaler.fit_transform(X_train[["age", "trestbps", "chol", "thalach", "oldpeak"]])
X_train.head()
```

H:\download\Anaconda\lib\site-packages\pandas\core\frame.py:3678: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

self[col] = igetitem(value, i)

```
Out[58]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
937	0.729167	1	2	0.534884	0.779851	0	1	0.614035	0	0.000000	2	0	2

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
953	0.333333	0	0	0.511628	0.410448	0	0	0.561404	1	0.045455	1	0	2
686	0.479167	1	0	0.395349	0.291045	1	1	0.596491	1	0.227273	1	0	0
791	0.520833	1	0	0.186047	0.421642	0	1	0.333333	1	0.636364	1	1	3
939	0.416667	0	1	0.465116	0.541045	0	1	0.649123	0	0.000000	1	0	2

In [59]: `X_test[["age", "trestbps", "chol", "thalach", "oldpeak"]] = scaler.transform(X_test[["age", "trestbps", "chol", "thalach", "oldpeak"]])`

H:\download\Anaconda\lib\site-packages\pandas\core\frame.py:3678: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

`self[col] = igetitem(value, i)`

Out[59]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
948	0.854167	1	0	0.593023	0.179104	0	1	0.324561	1	0.590909	0	0	3
546	0.854167	1	0	0.418605	0.731343	0	0	0.184211	0	0.545455	1	3	2
589	0.520833	1	0	0.325581	0.597015	0	0	0.245614	1	0.727273	1	2	2
681	0.625000	1	0	0.883721	0.746269	0	0	0.456140	1	0.772727	0	0	3
955	0.687500	1	2	0.418605	0.391791	0	1	0.508772	0	0.409091	1	3	3

In []:

In [60]:

```

from sklearn import model_selection
from sklearn.model_selection import cross_val_score
import xgboost as xgb
# function initializing baseline machine learning models
def GetBasedModel():
    basedModels = []
    basedModels.append(('LR_L2', LogisticRegression(penalty='l2')))
    basedModels.append(('KNN7', KNeighborsClassifier(7)))
    basedModels.append(('KNN9', KNeighborsClassifier(9)))
    basedModels.append(('CART', DecisionTreeClassifier()))
    basedModels.append(('NB', GaussianNB()))
    basedModels.append(('SVM Linear', SVC(kernel='linear', gamma='auto', probability=True)))
    basedModels.append(('SVM RBF', SVC(kernel='rbf', gamma='auto', probability=True)))
    basedModels.append(('RF_Ent100', RandomForestClassifier(criterion='entropy', n_estimators=100)))
    basedModels.append(('RF_Gini100', RandomForestClassifier(criterion='gini', n_estimators=100)))
    basedModels.append(('ET100', ExtraTreesClassifier(n_estimators=100)))
    basedModels.append(('XGB_100', xgb.XGBClassifier(n_estimators=100)))

    return basedModels

# function for performing 10-fold cross validation of all the baseline models

```

```
def BasedLine2(X_train, y_train, models):
    # Test options and evaluation metric
    num_folds = 10
    scoring = 'accuracy'
    seed = 7
    results = []
    names = []
    for name, model in models:
        kfold = model_selection.KFold(n_splits=10)
        cv_results = model_selection.cross_val_score(model, X_train, y_train, cv=kfold,
        results.append(cv_results)
        names.append(name)
        msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
        print(msg)

    return results, msg
```

```
In [61]: models = GetBasedModel()
names, results = BasedLine2(X_train, y_train, models)
```

```
LR_L2: 0.847848 (0.033783)
KNN7: 0.854114 (0.031305)
KNN9: 0.870491 (0.044680)
CART: 0.987405 (0.015962)
NB: 0.816377 (0.015779)
SVM Linear: 0.836519 (0.037641)
SVM RBF: 0.884320 (0.040066)
RF_Ent100: 0.983623 (0.018749)
RF_Gini100: 0.982358 (0.016137)
ET100: 0.984889 (0.016710)
XGB_100: 0.986155 (0.015355)
```

Random Forest Classifier (criterion = 'entropy')

```
In [62]: rf_ent = RandomForestClassifier(criterion='entropy', n_estimators=100)
rf_ent.fit(X_train, y_train)
y_pred_rfe = rf_ent.predict(X_test)
```

Draw the Confusion matrix

```
In [63]: cm=confusion_matrix(y_test, y_pred_rfe)
sns.heatmap(cm, annot=True)

TN = cm[0][0]
FN = cm[1][0]
TP = cm[1][1]
FP = cm[0][1]
acc= accuracy_score(y_test, y_pred_rfe)
roc=roc_auc_score(y_test, y_pred_rfe)
loss_log = log_loss(y_test, y_pred_rfe)
prec = precision_score(y_test, y_pred_rfe)
rec = recall_score(y_test, y_pred_rfe)
f1 = f1_score(y_test, y_pred_rfe)


mathew = matthews_corrcoef(y_test, y_pred_rfe)
```

```
model_results =pd.DataFrame([[ 'Random Forest',acc, prec,rec, f1,roc,mathew,loss_log]],
                             columns = [ 'Model', 'Accuracy','Precision', 'Recall_Score', 'F1 Score','ROC', 'mathew_corrcoef', 'Log Loss'])

model_results
```

```
Out[63]:
```

	Model	Accuracy	Precision	Recall_Score	F1 Score	ROC	mathew_corrcoef	Log Loss
0	Random Forest	1.0	1.0	1.0	1.0	1.0	1.0	9.992007e-16



Soft voting

```
In [64]: import xgboost as xgb
clf1=RandomForestClassifier(criterion='entropy',n_estimators=100)

clf2=DecisionTreeClassifier()
clf3=xgb.XGBClassifier(n_estimators= 1000)
clf4=ExtraTreesClassifier(n_estimators= 500)

clf5=GradientBoostingClassifier(n_estimators=100,max_features='sqrt')

eclf1 = VotingClassifier(estimators=[('rfe', clf1), ('decc', clf2), ('xgb', clf3),('ET',
                                     voting='soft', weights=[4,1,2,3,1])])
eclf1.fit(X_train,y_train)
y_pred_sv =eclf1.predict(X_test)
```

```
In [ ]:
```

```
In [65]: CM=confusion_matrix(y_test,y_pred_sv)
sns.heatmap(CM, annot=True)

TN = CM[0][0]
FN = CM[1][0]
TP = CM[1][1]
FP = CM[0][1]
loss_log = log_loss(y_test, y_pred_sv)
acc= accuracy_score(y_test, y_pred_sv)
roc=roc_auc_score(y_test, y_pred_sv)
prec = precision_score(y_test, y_pred_sv)
```



```

rec = recall_score(y_test, y_pred_sv)
f1 = f1_score(y_test, y_pred_sv)


mathew = matthews_corrcoef(y_test, y_pred_sv)
model_results = pd.DataFrame([[ 'Soft Voting', acc, prec, rec, f1, roc, loss_log, mathew ]],
                              columns = [ 'Model', 'Accuracy', 'Precision', 'Recall Score', 'F1 Score', 'ROC', 'Log_Loss', 'mathew_corrcoef' ])

model_results

```

Out[65]:

	Model	Accuracy	Precision	Recall Score	F1 Score	ROC	Log_Loss	mathew_corrcoef
0	Soft Voting	0.984925	0.972973	1.0	0.986301	0.983516	0.520697	0.969998



```

In [66]: rf_ent = RandomForestClassifier(criterion='entropy', n_estimators=100)
rf_ent.fit(X_train, y_train)
y_pred_rfe = rf_ent.predict(X_test)

```

```

In [67]: knn = KNeighborsClassifier(9)
knn.fit(X_train, y_train)
y_pred_knn = knn.predict(X_test)

```

```

In [68]: et_100 = ExtraTreesClassifier(n_estimators= 100)
et_100.fit(X_train, y_train)
y_pred_et100 = et_100.predict(X_test)

```

```

In [69]: import xgboost as xgb
xgb = xgb.XGBClassifier(n_estimators= 1000)
xgb.fit(X_train, y_train)
y_pred_xgb = xgb.predict(X_test)

```

```

In [70]: svc = SVC(kernel='linear', gamma='auto', probability=True)
svc.fit(X_train, y_train)
y_pred_svc = svc.predict(X_test)

```

```

In [71]: decc = DecisionTreeClassifier()
decc.fit(X_train, y_train)

```

```
y_pred_decc = decc.predict(X_test)
```

In [72]:

```
data = {
    'Random Forest Entropy': y_pred_rfe,
    'KNN2': y_pred_knn,
    'EXtra tree classifier': y_pred_et100,
    'XGB2': y_pred_xgb,
    'SVC2': y_pred_svc,
    'CART': y_pred_decc
}

models = pd.DataFrame(data)

for column in models:
    CM=confusion_matrix(y_test,models[column])

    TN = CM[0][0]
    FN = CM[1][0]
    TP = CM[1][1]
    FP = CM[0][1]
    loss_log = log_loss(y_test, models[column])
    acc= accuracy_score(y_test, models[column])
    roc=roc_auc_score(y_test, models[column])
    prec = precision_score(y_test, models[column])
    rec = recall_score(y_test, models[column])
    f1 = f1_score(y_test, models[column])

    mathew = matthews_corrcoef(y_test, models[column])
    results =pd.DataFrame([[column,acc, prec,rec, f1,roc, loss_log,mathew]],
        columns = ['Model', 'Accuracy', 'Precision', 'Recall Score', 'F1 Score', '
    model_results = model_results.append(results, ignore_index = True)

model_results
```

Out[72]:

	Model	Accuracy	Precision	Recall Score	F1 Score	ROC	Log_Loss	mathew_corrcoef
0	Soft Voting	0.984925	0.972973	1.000000	0.986301	0.983516	5.206971e-01	0.969998
1	Random Forest Entropy	1.000000	1.000000	1.000000	1.000000	1.000000	9.992007e-16	1.000000
2	KNN2	0.889447	0.898148	0.898148	0.898148	0.888635	3.818401e+00	0.777269
3	EXtra tree classifier	1.000000	1.000000	1.000000	1.000000	1.000000	9.992007e-16	1.000000
4	XGB2	0.984925	0.972973	1.000000	0.986301	0.983516	5.206971e-01	0.969998
5	SVC2	0.819095	0.846154	0.814815	0.830189	0.819495	6.248285e+00	0.637307
6	CART	0.969849	0.972222	0.972222	0.972222	0.969628	1.041382e+00	0.939255

In []: